

Resource Scheduling Problem in Hadoop

Project for Algorithm and Complexity Course

Department of Computer Science and Engineering,
Shanghai Jiao Tong University, Shanghai, China

Abstract. The course project focuses on resource scheduling problem in Hadoop. This document introduce the background and two versions of resource scheduling on single or multiple hosts. Also, it lists all the tasks and requirements for each student group. Please read this document carefully and complete the corresponding tasks.

Keywords: Distributed Computing System, Resource Scheduling, Hadoop

1 Background and Motivation

Hadoop is an open-source software framework for storing data and running applications on clusters of commodity hardware. In Hadoop, the data of a running job can be divided into several data blocks, stored on different hosts. Each host has one or multiple CPU cores to process data blocks. In this project, our goal is to achieve effective parallel computing. That is to say, given a set of jobs with massive data, we need to design a resource scheduling algorithm to minimize the overall executing time of all jobs.

2 A Simplified Version with Single Host

First, let us consider a simple case with a single host storing all data blocks. Please read the assumptions, specifications, symbol notations, constraints, and explanations in Subsection 2.1 carefully and then complete the tasks mentioned in Subsection 2.2.

2.1 Problem Formulation and Explanation

To simplify the problem, we give the following assumptions and specifications.

1. There are n jobs that need to be processed by a single host, which has m CPU cores of the same computing capability. Let J be the set of jobs, and C be the set of cores for the host, where $J = \{job_0, job_1, \dots, job_{n-1}\}$, and $C = \{c_0, c_1, \dots, c_{m-1}\}$ (The labels of variables start from 0 because we use C/C++ source codes in the following tasks).
2. We treat data block as the smallest indivisible unit in our project, while a job can be divided into multiple data blocks with different sizes for storage and computing. Assume job_i is split into n_i data blocks, denoted by $B^i = \{b_0^i, b_1^i, \dots, b_{n_i-1}^i\}$. For block b_k^i of job_i , define its size as $size(b_k^i)$.
3. Assume job_i is assigned to e_i cores for processing, and naturally $e_i \leq n_i$. That is to say, one core can process multiple data blocks of one job sequentially. Let $B_j^i \subseteq B^i$ denote the set of data blocks of job_i allocated to core c_j , and $B_j^i \cap B_{j'}^i = \emptyset$ if $j \neq j'$ (they should be disjointed).
4. For job_i , the processing speed of its data blocks is s_i when job_i is assigned to a single core. However, when multiple cores process job_i in parallel, the computing speed of each core all decays because of some complicated interactions. We formulate such speed decay effect caused by multi-core computation as a coefficient function $g(\cdot)$ with respect to core number e_i , as described in Equation (1):

$$g(e_i) = 1.00 - \alpha \cdot (e_i - 1), \quad \text{for } 1 \leq e_i \leq 10, \quad (1)$$

where α is a decay factor satisfying $0 < \alpha < 1$, and usually the number of cores for processing a single job is no more than 10. Then, the speed of each core can be rewritten as $s_i \cdot g(e_i)$ for job_i respectively. (Note that although the speed of each core decays, the overall processing time using e_i cores in parallel should be faster than that of using just one core. Otherwise we do not need to implement parallel computing. Thus the setting of α should guarantee this principle.)

Correspondingly, the processing time tp_j^i of core c_j for job_i can be expressed as Equation (2):

$$tp_j^i = \frac{\sum_{b_k^i \in B_j^i} size(b_k^i)}{s_i \cdot g(e_i)}. \quad (2)$$

5. For consistency issues, if we assign a job to multiple cores, all cores must start processing data blocks at the same time. If one or several cores are occupied by other affairs, then all other cores should wait for a synchronous start, and keep idle. It means that the processing of job_i for every core should all start at time t_i , whereas their processing duration might be different. Let tf_j^i be the finishing time of core c_j for job_i , which is calculated by Equation (3):

$$tf_j^i = t_i + tp_j^i. \quad (3)$$

However, the occupied cores of job_i are released synchronously when the computing process of the last data block is finished. Thus the finishing time $tf(job_i)$ of job_i is given as Equation (4):

$$tf(job_i) = \max_{c_j \in C} tf_j^i, \text{ for } c_j \in C. \quad (4)$$

Please keep these assumptions and specifications in mind and finish the following tasks.

2.2 Task 1: Resource Scheduling for Single Host

Based on the descriptions in Subsection 2.1, your work is to design a resource scheduling algorithm to minimize the overall finishing time of all jobs, whose objective function is shown as:

$$\min \max_{job_i} tf(job_i), \text{ for } job_i \in J.$$

Figure 1 illustrates a toy example of scheduling resources among three jobs job_0, job_1, job_2 on a single host, which has four cores c_0, c_1, c_2, c_3 . All blocks of job_0, job_1 , and job_2 are stored on this host, with different sizes measured by megabyte (MB). The goal is to find a scheduling strategy assigning data blocks to suitable cores so that the overall finishing time of three jobs is minimized.

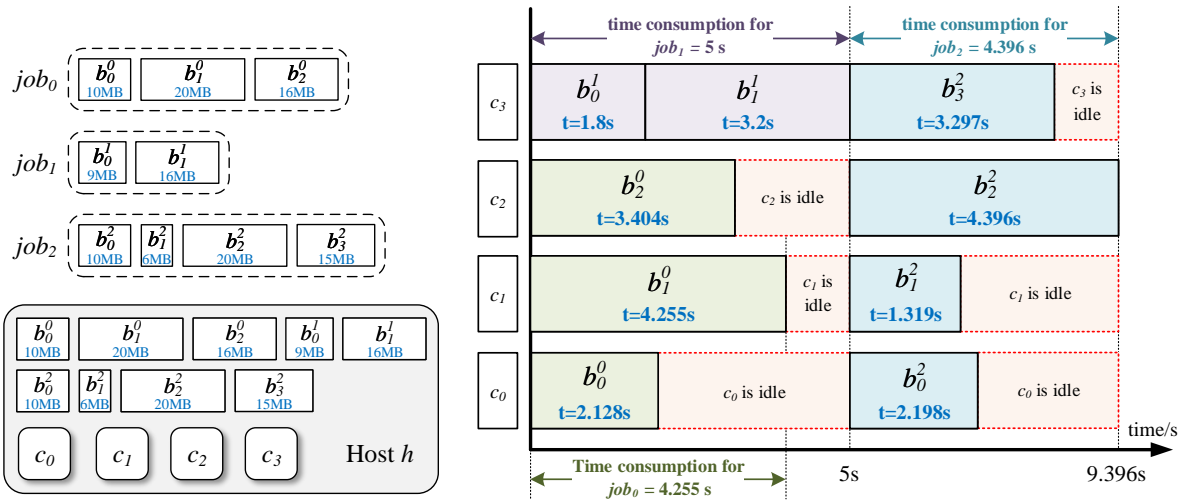


Fig. 1. A toy example of scheduling resources for a single host

Fig. 2. Time consumption of a feasible solution for the example in Figure 1

Remark. Assume the sizes of three blocks of job_0 are respectively $size(b_0^0) = 10$ MB, $size(b_1^0) = 20$ MB and $size(b_2^0) = 16$ MB, while those of job_1 are $size(b_0^1) = 9$ MB and $size(b_1^1) = 16$ MB. The size of each block in job_2 is $size(b_0^2) = 10$ MB, $size(b_1^2) = 6$ MB, $size(b_2^2) = 20$ MB and $size(b_3^2) = 15$ MB, respectively. Moreover, suppose the computing speed is $s_i = 5$ MB/s for $i = 0, 1, 2$, and $\alpha = 0.03$ to compute the decay coefficient $g(\cdot)$ in Equation (1).

Here we provide a feasible solution (as shown in Figure 2) for the above setting in Figure 1, in which blocks of job_0 are assigned to three cores c_0, c_1, c_2 , and blocks of job_1 are assigned to the last core c_3 .

Now, we can compute the processing time of core c_0, c_1, c_2 for job_0 as Equation (5) respectively:

$$\begin{cases} tp_0^0 = size(b_0^0)/(s_0 \cdot g(e_0)) = 10/(5 \times (1.00 - 0.03 \times (3 - 1))) = 2.128s, \\ tp_1^0 = size(b_1^0)/(s_0 \cdot g(e_0)) = 20/4.7 = 4.255s, \\ tp_2^0 = size(b_2^0)/(s_0 \cdot g(e_0)) = 16/4.7 = 3.404s. \end{cases} \quad (5)$$

Moreover, the processing of job_0 starts at 0s, so the finishing time of job_0 is

$$tf(job_0) = 0 + \max\{tp_0^0, tp_1^0, tp_2^0\} = \max\{2.128, 4.255, 3.404\} = 4.255s.$$

The processing time of blocks b_0^1 and b_1^1 with the single core c_3 is $t_0^1 = 9/5 = 1.8s$ and $t_1^1 = 16/5 = 3.2s$, respectively. Then the finishing time of job_1 is $tf(job_1) = 1.8 + 3.2 = 5s$.

We assign job_2 to four cores, each of which should start after job_1 releases the occupied core. Obviously, the processing of job_2 starts at 5s. The processing time of each block of job_2 is:

$$\begin{cases} t_0^2 = size(b_0^2)/(s_2 \cdot g(e_2)) = 10/(5 \times (1.00 - 0.03 \times (4 - 1))) = 2.198s, \\ t_1^2 = size(b_1^2)/(s_2 \cdot g(e_2)) = 6/4.55 = 1.319s, \\ t_2^2 = size(b_2^2)/(s_2 \cdot g(e_2)) = 20/4.55 = 4.396s, \\ t_3^2 = size(b_3^2)/(s_2 \cdot g(e_2)) = 15/4.55 = 3.297s. \end{cases}$$

Then we can get that the finishing time of job_2 as

$$tf(job_2) = 5 + \max\{2.198, 1.319, 4.396, 3.297\} = 9.396s.$$

Thus, the overall finishing time of the three jobs is

$$\max\{tf(job_0), tf(job_1), tf(job_2)\} = 9.396s.$$

It is observed that for a multi-core job, its time consumption is determined by the last completed block, while for the whole system, the total finishing time is decided by the last finished job. However, Figure 2 is just a feasible solution and not necessarily optimal, since cores c_0, c_1, c_2 , and c_3 are all idle for a while during the whole process, which is a waste of resources.

Based on the above explanations and examples, please finish the following tasks:

1. Formalize the resource scheduling problem for a single host as a programming pattern with objective function and constraints. You cannot rename the variables that have been defined in the project, whereas you are free to introduce other new variables, and please define your variables clearly.
2. Please design an algorithm to solve this problem efficiently. First, describe your idea in detail, and then provide the corresponding pseudocode. Also, please discuss the time complexity of your design.
3. Verify your algorithm using the attached test data in “task1_case1.txt” under *input* file folder and save your result in the .txt file, named as “task1_case1_TeamNumber.txt” (e.g., task1_case1_06.txt). The input/output format is fixed in the reference codes. Optionally, visualize your result while the visual format is not limited. For example, you can plot a figure from the perspective of cores on the host. The test data and reference codes are also released on GitHub: [Resource Scheduling Problem](#).

3 A Comprehensive Version among Multiple Hosts

In this section, we consider a more complex situation, where we need to schedule resources among multiple hosts. Now the data transmission process between pairwise hosts should be taken into consideration. The data blocks of jobs could be initially stored on different hosts, but one data block can only be initially stored on one specified host. If data block b_k^i and its assigned computing core c_j are not on the same host, b_k^i will be transmitted to the host containing c_j (Here we assume that the bandwidth between hosts is sufficient for data transmission). The transmission process will influence the finishing time of jobs, further affecting the resource scheduling process.

3.1 Problem Formulation and Explanation

Besides the descriptions and specifications of Task 1, here are more notations and explanations.

1. Assume we have q hosts $H = \{h_0, h_1, \dots, h_{q-1}\}$, and host h_l has m_l cores (may have different number of cores). Let C_l be the set of cores on host h_l , $C_l = \{c_0^l, c_1^l, \dots, c_{m_l-1}^l\}$. Easy to see, $\sum_{l=0}^{q-1} m_l = m$.
2. If core c_j^l on host h_l computes a data block b_k^i of job_i which is initially stored on another host $h_{l'}$, then b_k^i needs to be transmitted from $h_{l'}$ to h_l at a transmission speed s_t (this speed is fixed in our system). An example is shown in Figure 3, where hosts h_0 and h_1 both have two cores, and many jobs need to be processed. Core c_0^1 on host h_1 is assigned to compute the data block b_2^1 of job_1 , which is initially stored on host h_0 . In this case, b_2^1 needs to be transmitted from h_0 to h_1 at a transmission speed s_t first, and then be computed by c_0^1 . Whenever b_2^1 starts transmission, other cores can work in parallel to process job_1 .

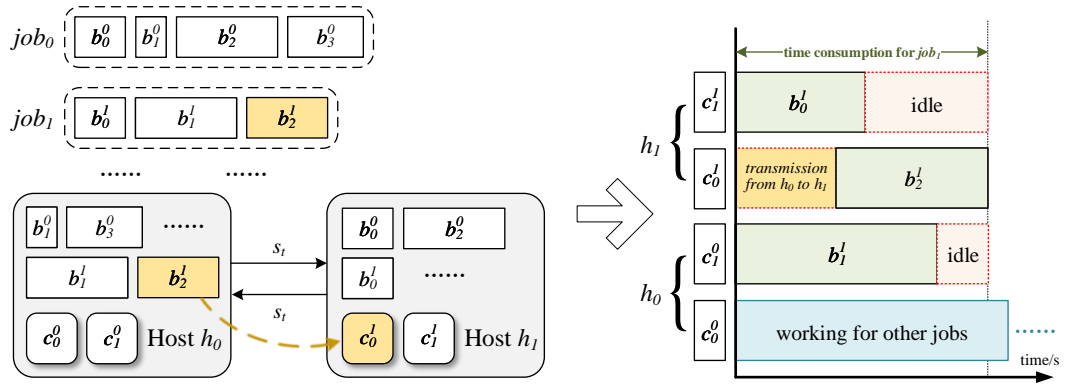


Fig. 3. An example of data transmission between 2 hosts

3. Any core cannot call for data transmission when it is calculating other data blocks. Likewise, a core cannot start computing any data block until this block is ready, i.e. initially on the same host or transmitted from a remote host to the local host. For example, the core c_0^1 on host h_1 must wait for the data transmission of block b_2^1 from host h_0 to h_1 , and then start computation. What is more, the transmission time of b_2^1 from h_0 to h_1 affects the finishing time of job_0 , further affecting the finishing time of the whole system.

For core c_j^l on host h_l , let \tilde{B}_{lj}^i be the set of data blocks of job_i allocated to c_j^l but not initially stored on host h_l . All the data blocks in \tilde{B}_{lj}^i need to be transmitted to host h_l before computing. Let B_{lj}^i be the set of data blocks of job_i allocated to core c_j^l . Then, the processing time tp_{lj}^i of core c_j^l for job_i can be reformulated as Equation (6):

$$tp_{lj}^i = \frac{\sum_{b_k^i \in \tilde{B}_{lj}^i} size(b_k^i)}{s_t} + \frac{\sum_{b_k^i \in B_{lj}^i} size(b_k^i)}{s_i \cdot g(e_i)}. \quad (6)$$

4. If the processing of job_i starts at time t_i , then the finishing time of core c_j^l for job_i is

$$tf_{lj}^i = t_i + tp_{lj}^i.$$

Then the finishing time $tf(job_i)$ of job_i is formulated as:

$$tf(job_i) = \max_{c_j^l \in C} tf_{lj}^i, \text{ for } c_j^l \in C.$$

Please keep these assumptions and specifications in mind and finish the following tasks.

3.2 Task 2: Resource Scheduling among Multiple Hosts

Similarly, the aim of Task 2 is to design a resource scheduling algorithm among multiple hosts to minimize the overall finishing time, which is formulated as:

$$\min \max_{job_i} tf(job_i), \text{ for } job_i \in J.$$

Figure 4 shows a toy example of scheduling resources among 3 hosts. We have 4 different jobs, each of which consists of blocks with different sizes. For example, job_0 has three blocks, b_0^0 , b_1^0 , and b_2^0 , stored on host h_0 . Other jobs are stored on different hosts. These data blocks are assigned to available cores on the three hosts, each with **two** cores. Then this scheduling should consider the data transmission process if the data block and its allocated core are not on the same host. Our goal is to find a scheduling strategy to allocate data blocks to appropriate cores so that the overall finishing time of four jobs is minimized.

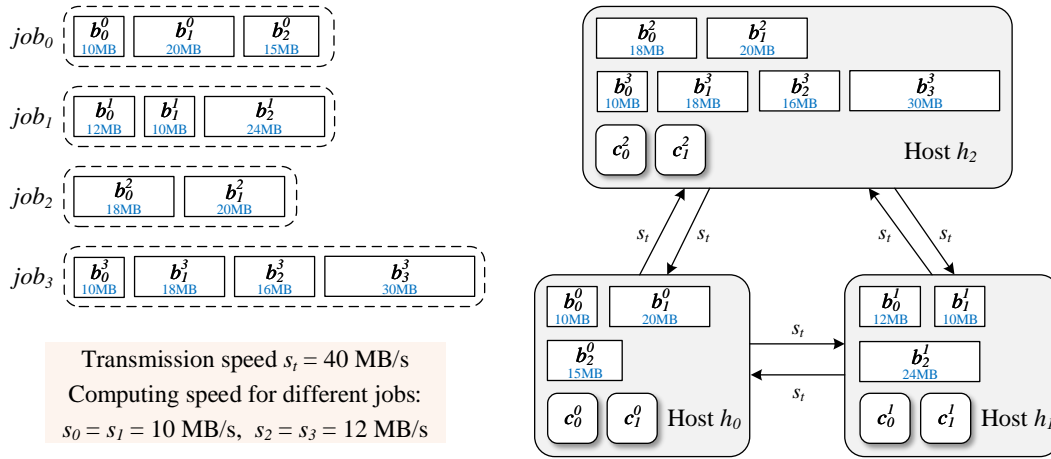


Fig. 4. A toy example of scheduling resources among 3 hosts

Remark. Assume the sizes of three blocks of job_0 are respectively set as $size(b_0^0) = 10$ MB, $size(b_1^0) = 20$ MB and $size(b_2^0) = 15$ MB, while those of other jobs are listed in the left hand of Figure 4. Set decay factor $\alpha = 0.1$, and then the computing decaying coefficient is $g(e_i) = 1 - 0.1 \times (e_i - 1)$. Besides, we set the transmission speed as $s_t = 40$ MB/s. Assume that job_0 and job_1 have the same computing speed, which is $s_0 = s_1 = 10$ MB/s. Similarly, the computing speed of job_2 and job_3 is $s_2 = s_3 = 12$ MB/s. Under the above settings, we give a feasible solution for the example in Figure 4, as shown in Figure 5.

In the initialization phase, we know that blocks of job_0 are on the host h_0 , blocks of job_1 are on h_1 , blocks of job_2 and job_3 are on h_2 . Thus we choose to compute job_0 , job_1 , and job_2 on their own local host, each with two cores.

Firstly, the calculation of time consumption is similar to the example of Task 1. For instance, the time consumption of block b_0^2 , also the processing time of core c_0^2 for job_2 , is

$$tp_{21}^2 = \frac{size(b_0^2)}{s_2 \times g(e_2)} = \frac{18}{12 \times g(2)} = \frac{18}{12 \times 0.9} = 1.667s. \quad (7)$$

As shown in Equation (7), we can compute the time consumption of the rest blocks in job_0 , job_1 and job_2 . One job can be allocated to multiple cores on different hosts and each block of this job can be computed by only one core. We allocate the data blocks of job_2 to cores c_0^2 and c_1^2 for computing. In detail, c_0^2 computes b_0^2 and c_1^2 computes b_1^2 . Cores c_0^2 and c_1^2 must be released simultaneously when the computing of b_1^2 is finished. Thus, the finishing time of job_2 in this stage is $1.852s$, while c_0^2 and c_1^2 stay idle to wait for new task allocation.

In the same way, job_0 , job_1 , and job_2 have been allocated to certain cores to finish computing. As shown in Figure 5, 4 cores on h_1 and h_2 stay in idle state after finishing the computation of c_1^1 . We consider that allocating job_3 to 4 cores on h_1 and h_2 is feasible.

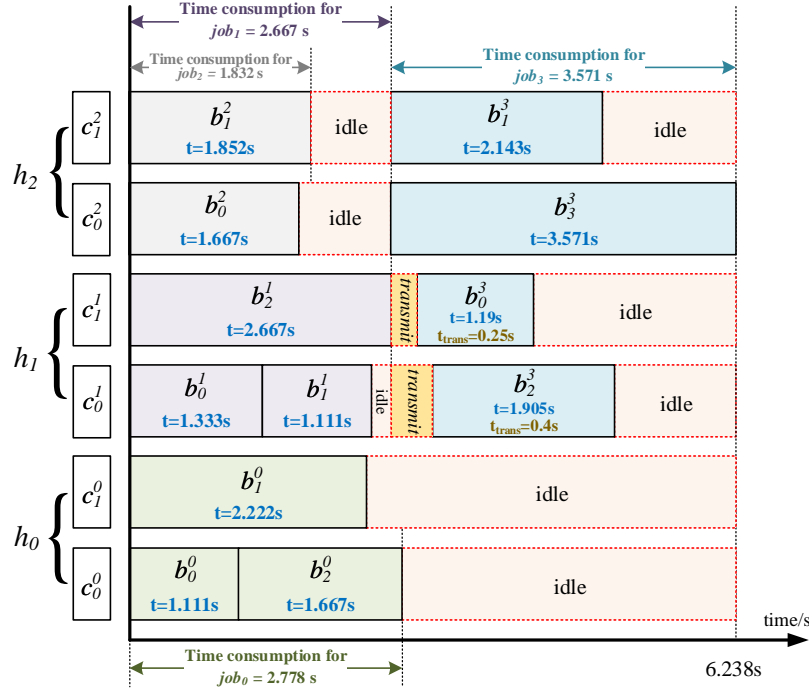


Fig. 5. Time consumption of a feasible solution for the example in Figure 4.

When the system allocates job_3 to 4 cores and starts processing it, it is **necessary** to transmit the blocks of job_3 from host h_2 to host h_1 . According to the actual condition, the time of job computing starting is the time of transmission starting. For instance, 4 blocks of job_3 are allocated to 4 cores for computing. We can compute b_1^3 and b_3^3 at local host directly, rather transmit b_0^3 and b_2^3 to host h_1 . The transmission time for b_0^3 can be calculated as $\frac{size(b_0^3)}{s_t} = \frac{10}{40} = 0.25s$.

Consequently, the time consumption of b_0^3 , also the processing time of core c_1^1 for job_3 , could be computed as Equation (8).

$$tp_{11}^3 = 0.25 + \frac{size(b_0^3)}{s_3 \times g(e_3)} = 0.25 + \frac{10}{12 \times g(4)} = 0.25 + \frac{10}{12 \times 0.7} = 1.440s. \quad (8)$$

It is the same as the above that some cores could be in idle state when finishing block computing. Thus, the processing time of job_3 in this stage is the time consumption of b_3^3 , which is 3.571s. The whole time consumption of this example is the latest finishing time of all jobs, which is calculated to be 6.238s. Figure 5 shows the time consumption of this solution and the red dotted box expresses the idle state of cores during the scheduling process.

Obviously, the time consumption of this solution could utilize resources as much as possible, which means that the data transmission can affect the final overall finishing time. However, what is remarkable is that the solution given in Figure 5 may not be optimal for the example in Figure 4. Because idle state indicates the resource waste and this solution can still be optimized.

Based on the above discussions and examples, please finish the following tasks:

1. Formalize the resource scheduling problem among multiple hosts as a programming pattern with objective function and constraints. Justify whether it is an NP-Complete problem.
2. Design an algorithm to solve this problem. First, describe your idea in detail, and then provide the corresponding pseudo code. Also, please discuss the time complexity of your design.
3. Verify your algorithm using test data "task2_case1.txt" in *input* file folder and save your result in the .txt file, named as "task2_case1_TeamNumber.txt" (e.g., "task2_case1_06.txt"). The input/output format is fixed in the reference codes. Optionally, visualize your result in an unlimited format. For example, you can plot a figure from the perspective of cores on the host. The test data and reference codes are also released on GitHub: [Resource Scheduling Problem](#).

4 Report Requirements

You need to submit a report for this project, with the following requirements:

1. Your report should include the title, the author names, IDs, email addresses, the page header, the page numbers, your results and discussions for the tasks, figures for your simulations, tables for discussions and comparisons, with the corresponding figure titles and table titles.
2. Your report should be in English only, with a clear structure, divided by sections, and may contain organizational architecture such as items, definitions, or discussions.
3. Please include Reference section and Acknowledgment section. You may also include your feelings, suggestions, and comments in the acknowledgment section.
4. Please define your variables clearly and add them into the symbol table in Appendix.
5. Please create a folder named "Project-TeamNumber" which contains related materials such as report "Project-TeamNumber.pdf", latex source "Project-TeamNumber.tex", the executable file "Project-TeamNumber.exe" (if you have), the data output folder "Project-Outputs-TeamNumber", the figure folder "Project-Figures-TeamNumber", and other code file folder "Project-Codes-TeamNumber". Then compress the home folder "Project-TeamNumber" into a "Project-TeamNumber.zip" package. Note that TeamNumber should be two-digit number, e.g., "Project-06.zip" conforms to the rule.

Acknowledgements

This problem is motivated from a real-world cloud service corporation. It is formulated by Prof. Xiaofeng Gao (gao-xf@cs.sjtu.edu.cn) from Department of Computer Science and Engineering at Shanghai Jiao Tong University. Yue Ma scribed and modified the project. Yangyang Bao provided the example in Task 1. Zhen Sun provided the example in Task 2. Wanghua Shi provided the test data and source code. Jiale Zhang, Tianyao Shi helped on proofreading and provided many suggestions.

Appendix

Table 1. Symbols and Definitions

| Symbols | Definitions |
|------------------------|-------------------------------------------------------------------------------------------------------------|
| n | The number of jobs |
| m | The number of cores |
| q | The number of hosts |
| job_i, J | job_i is the i -th job. The job set is $J = \{job_0, \dots, job_{n-1}\}$. |
| h_l, H | h_l is the l -th host. The host set is $H = \{h_0, \dots, h_{q-1}\}$. |
| m_l | The number of cores on host h_l |
| c_j^l, C_l | c_j^l is the j -th core on host h_l . C_l is the set of cores on host h_l . |
| C | The set of cores. $C = \{c_0, \dots, c_{m-1}\}$ for single-host. $C = \cup_{l=0}^{q-1} C_l$ for multi-host. |
| b_k^i | The block of job_i whose id is k |
| B_j^i | The set of data blocks of job_i allocated to core c_j |
| B^i | The set of data blocks of job_i |
| B_{lj}^i | The set of data blocks of job_i allocated to core c_j^l |
| \tilde{B}_{lj}^i | The set of data blocks of job_i allocated to core c_j^l but not initially stored on h_l |
| $size(\cdot)$ | The size function of data block |
| $g(\cdot)$ | The computing decaying coefficient caused by multi-core effect |
| s_i | The computing speed of job_i by a single core |
| s_t | The transmission speed of data |
| e_i | The number of cores processing job_i |
| t_i | The time to start processing job_i |
| tp_j^i, tf_j^i | The processing time / finishing time of core c_j for job_i |
| tp_{lj}^i, tf_{lj}^i | The processing time / finishing time of core c_j^l for job_i |
| $tf(job_i)$ | The finishing time of job_i |